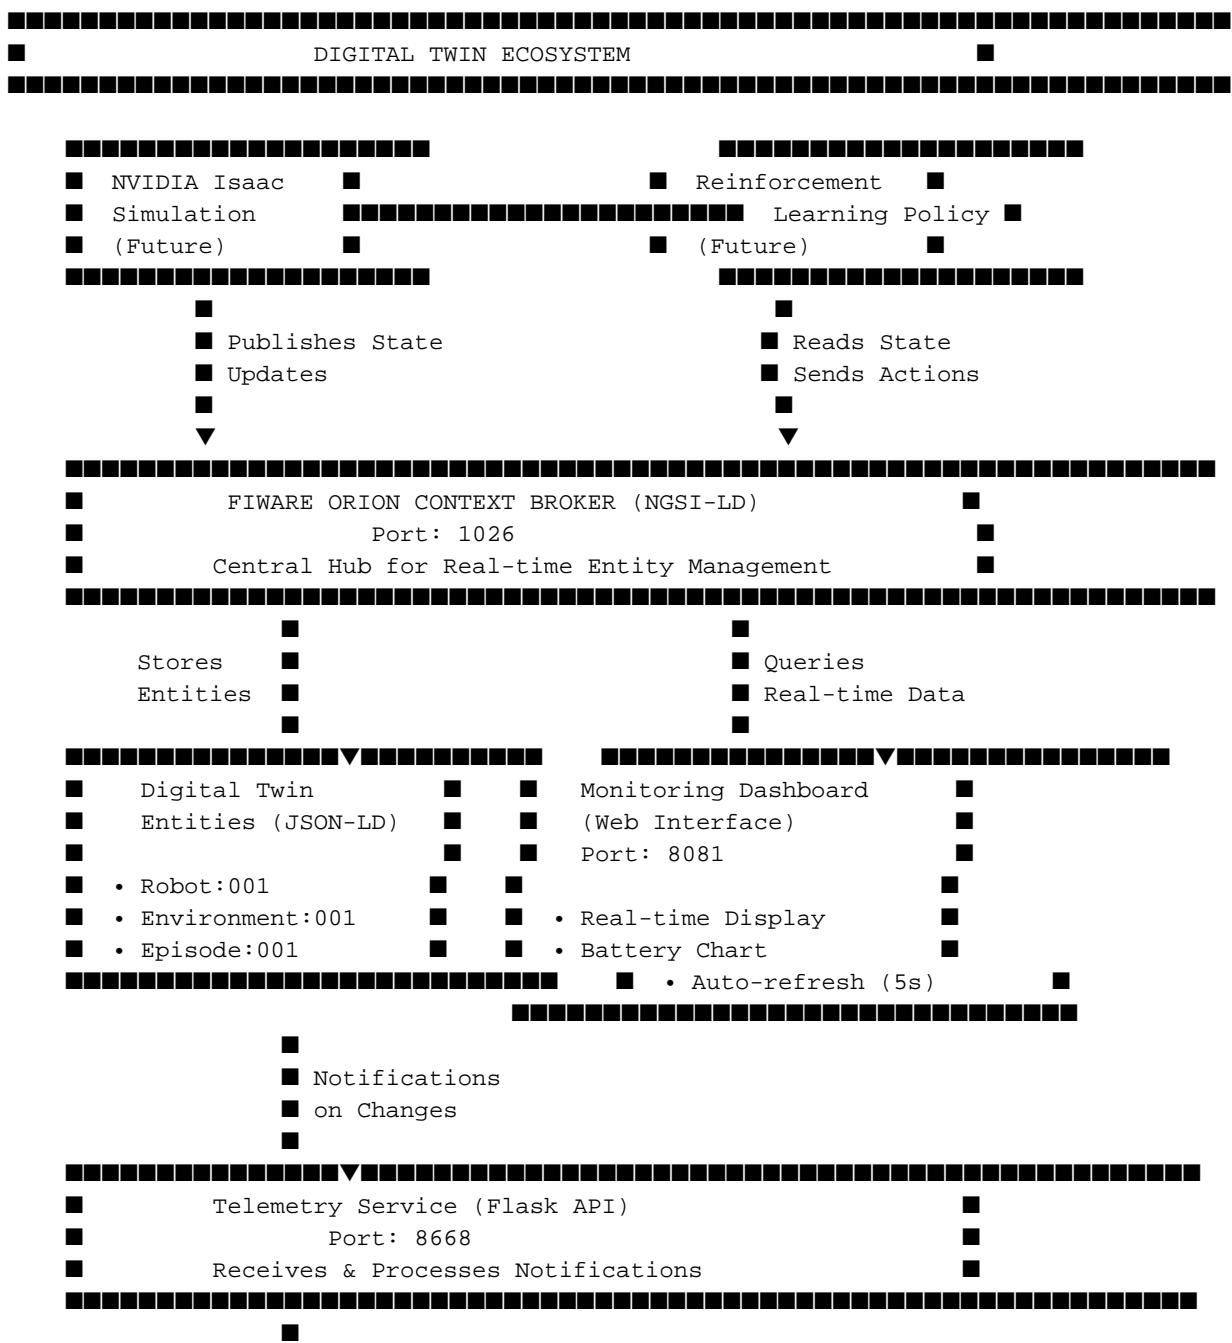


Digital Twin Architecture Documentation

System Overview

This document explains the complete Digital Twin architecture for Robot Reinforcement Learning in Logistics, detailing each component's role and how they interact.

Architecture Diagram





Component Details

1. FIWARE Orion Context Broker (NGSI-LD)

Role: Central nervous system of the digital twin architecture

Technology:

- FIWARE Orion-LD v1.5.1
- NGSI-LD (Next Generation Service Interface - Linked Data)
- MongoDB 4.4 backend

Responsibilities:

Entity Management: Stores and manages digital twin entities

- Robot entities (position, velocity, battery, sensors)
- Environment entities (temperature, lighting, obstacles)
- Episode entities (training metrics, rewards, steps)

Real-time Updates: Provides instant access to current state

- RESTful API endpoints for CRUD operations
- NGSI-LD format ensures semantic interoperability
- JSON-LD context for standardized vocabulary

Subscription System: Notifies subscribers when entities change

- Publishes notifications to telemetry service
- Supports attribute-level filtering
- Enables event-driven architecture

CORS Support: Allows dashboard to fetch data

- Configured with `-corsOrigin __ALL`
- Enables cross-origin requests from web browser

Port: 1026

Key Endpoints:

- `GET /version` - Check Orion version and status
- `POST /ngsi-ld/v1/entities` - Create new entity
- `GET /ngsi-ld/v1/entities/{id}` - Retrieve entity
- `PATCH /ngsi-ld/v1/entities/{id}/attrs` - Update attributes
- `POST /v2/subscriptions` - Create subscription for notifications

Data Flow:

```

Isaac Sim → Updates Robot State → Orion Context Broker
RL Policy → Reads State from Orion → Computes Action → Updates Orion
Dashboard → Queries Orion → Displays Real-time Data

```

2. Digital Twin Entities (JSON-LD)

Role: Semantic representation of physical and virtual components

Technology:

- NGSI-LD data model
- JSON-LD with custom context
- Standardized vocabulary

Entity Types:

Robot Entity (<urn:nsgsi-ld:Robot:001>)

Purpose: Digital representation of the physical/simulated robot

Attributes:

- `position` (GeoProperty): 3D coordinates [x, y, z]
- `velocity` (Property): Speed vector {vx, vy, vz}
- `orientation` (Property): Quaternion {x, y, z, w}
- `battery` (Property): Charge level (0-100%)
- `status` (Property): Operational state (idle, moving, charging, error)
- `sensorData` (Property):
 - IMU: Acceleration and gyroscope readings
 - Lidar: Point cloud data
 - Camera: Image status

Use Case: Isaac Sim publishes robot state updates; RL policy reads state to decide actions

Environment Entity (<urn:nsgsi-ld:Environment:001>)

Purpose: Represents simulation environment conditions

Attributes:

- `temperature` (Property): Ambient temperature
- `lighting` (Property): Light configuration
- `obstacles` (Property): List of obstacles in scene
- `dimensions` (Property): Environment size

Use Case: Defines training environment parameters for RL agent

Episode Entity (<urn:nsgsi-ld:Episode:001>)

Purpose: Tracks reinforcement learning training progress

Attributes:

- `episodeNumber` (Property): Current training episode
- `reward` (Property): Accumulated reward
- `steps` (Property): Number of timesteps
- `metrics` (Property): Performance indicators
 - Success rate
 - Collision count
 - Distance traveled

Use Case: Monitors training progress and performance metrics

JSON-LD Context:

```
{  
  "@context": [  
    "https://uri.etsi.org/ngsi-ld/v1/ngsi-ld-core-context.jsonld",  
    {  
      "Robot": "https://example.org/digitaltwin/Robot",  
      "position": "https://example.org/digitaltwin/position",  
      "velocity": "https://example.org/digitaltwin/velocity"  
    }  
  ]  
}
```

3. Telemetry Service (Custom Flask Application)

Role: Historical data collection and storage bridge

Technology:

- Python Flask web framework
- Flask-CORS for cross-origin support
- psycopg2 for PostgreSQL connection

Responsibilities:

Notification Receiver: Listens for Context Broker notifications

- Endpoint: `POST /v2/notify`
- Receives entity change notifications via subscriptions
- Processes NGSI-LD attribute updates

Data Transformation: Converts real-time data to time-series format

- Extracts entity ID, type, attribute name, and value
- Adds timestamp for temporal ordering
- Stores metadata (dateCreated, dateModified)

Database Writer: Persists data to TimescaleDB

- Batch inserts for efficiency
- Uses hypertable for time-series optimization
- Indexes on entity_id and time for fast queries

Query API: Provides historical data access

- Endpoint: `GET /v2/entities/{id}?lastN=X`
- Returns last N records for an entity
- Supports time-range filtering

Health Monitoring: Status endpoint for system checks

- Endpoint: `GET /health`
- Returns database connection status

Port: 8668

Docker Configuration:

```
telemetry-service:  
  build: ./telemetry_store  
  ports: ["8668:8668"]  
  environment:  
    - DB_HOST=timescaledb  
    - DB_PORT=5432
```

Subscription Example:

```
{  
  "description": "Notify Telemetry Store of Robot changes",  
  "subject": {  
    "entities": [{ "idPattern": ".*", "type": "Robot" }],  
    "condition": { "attrs": [ "battery", "position", "status" ] }  
  },  
  "notification": {  
    "http": { "url": "http://telemetry-service:8668/v2/notify" },  
    "attrs": [ "battery", "position", "status" ]  
  }  
}
```

Data Flow:

```
Context Broker → Subscription Triggers → POST /v2/notify →  
Extract Attributes → Format for TimescaleDB → INSERT INTO telemetry_data
```

4. TimescaleDB (*Time-series Database*)

Role: Long-term storage for historical telemetry data

Technology:

- PostgreSQL 15 with TimescaleDB extension
- Hypertable optimization for time-series data
- Automatic partitioning by time

Responsibilities:

Time-series Storage: Efficiently stores timestamped data

- Automatic chunking by time intervals
- Compression for older data
- Fast time-range queries

Data Retention: Manages historical data lifecycle

- Configurable retention policies
- Automatic data aggregation
- Space optimization

Analytics Support: Enables historical analysis

- Aggregate functions (AVG, MIN, MAX, STDDEV)
- Time-bucket queries
- Trend analysis

Port: 5432

Database Schema:

```
CREATE TABLE telemetry_data (  
    id SERIAL,  
    time TIMESTAMPTZ NOT NULL,          -- When data was recorded  
    entity_id TEXT NOT NULL,            -- Which entity (Robot001)  
    entity_type TEXT NOT NULL,          -- Type of entity (Robot)  
    attribute_name TEXT NOT NULL,        -- What changed (battery)  
    attribute_value JSONB NOT NULL,       -- New value (85)  
    metadata JSONB,                     -- Additional context  
    PRIMARY KEY (id, time)  
);  
  
-- Convert to hypertable for time-series optimization  
SELECT create_hypertable('telemetry_data', 'time');  
  
-- Indexes for fast queries  
CREATE INDEX idx_entity_id ON telemetry_data (entity_id, time DESC);  
CREATE INDEX idx_entity_type ON telemetry_data (entity_type, time DESC);
```

Example Query:

```
-- Get last 10 battery readings for Robot001  
SELECT time, attribute_value  
FROM telemetry_data  
WHERE entity_id = 'Robot001'  
    AND attribute_name = 'battery'  
ORDER BY time DESC  
LIMIT 10;
```

Benefits:

- 10-100x faster than regular PostgreSQL for time-series queries
- Automatic data compression
- Native support for time-based partitioning
- SQL compatibility

5. Monitoring Dashboard (Web Interface)

Role: Real-time visualization and monitoring interface

Technology:

- Pure HTML/CSS/JavaScript
- Chart.js for data visualization
- Fetch API for REST calls
- Responsive design with CSS Grid

Features:****Real-time Display**:**

- Robot status (name, operational state)
- Battery level with visual bar indicator
- 3D position coordinates (X, Y, Z)
- Velocity vector (vX, vY, vZ)
- Sensor readings (IMU, Lidar, Camera)

****Training Metrics**:**

- Current episode number
- Accumulated reward
- Step count
- Performance indicators

****Environment Info**:**

- Temperature
- Lighting configuration
- Environment dimensions

****Historical Charts**:**

- Battery level over time (Chart.js)
- Last 20 data points from telemetry
- Interactive visualization

****Auto-refresh**:**

- Updates every 5 seconds
- Checks service health status
- Non-blocking async calls

Port: 8081 (HTTP server)

API Integration:

```
// Fetch robot data from Context Broker
const robot = await fetch(
  "http://localhost:1026/ngsi-ld/v1/entities/urn:ngsi-ld:Robot:001",
  { headers: { Accept: "application/ld+json" } }
);

// Fetch battery history from Telemetry Service
const history = await fetch(
  "http://localhost:8668/v2/entities/Robot001?lastN=20"
)
```

) ;

CORS Configuration:

- Orion: Returns `Access-Control-Allow-Origin: *`
- Telemetry Service: Uses flask-cors library
- Enables browser to fetch from different ports

User Interface:

- Purple gradient background
- Card-based layout
- Color-coded status badges
- Responsive grid system
- Battery indicator (green/yellow/red)

Data Flow Scenarios

Scenario 1: Robot State Update

1. Isaac Sim detects robot moved to new position
 - > Publishes update to Orion Context Broker
2. Orion receives PATCH request
 - > Updates Robot:001 entity attributes
 - > Stores in MongoDB
3. Orion checks active subscriptions
 - > Finds telemetry subscription for Robot entities
 - > Sends POST /v2/notify to telemetry-service:8668
4. Telemetry Service receives notification
 - > Extracts entity_id, attributes, timestamp
 - > Inserts into TimescaleDB telemetry_data table
5. Dashboard auto-refresh triggers (every 5s)
 - > Fetches latest state from Orion
 - > Fetches battery history from Telemetry Service
 - > Updates UI with new data

Scenario 2: Training Episode Completion

1. RL Policy completes episode
 - > Calculates final reward and metrics
 - > Updates Episode:001 entity in Orion
2. Orion notifies Telemetry Service
 - > Episode metrics stored in TimescaleDB
3. Dashboard displays
 - > Current episode number
 - > Total reward
 - > Success rate from metrics
4. Data available for analysis
 - > Query historical episode performance
 - > Track learning progress over time

Scenario 3: Historical Data Query

1. User wants to analyze battery usage patterns
 ■■> Opens dashboard battery chart
 2. Dashboard queries Telemetry Service
 ■■> GET /v2/entities/Robot001?lastN=20
 3. Telemetry Service queries TimescaleDB
 ■■> SELECT * FROM telemetry_data
 WHERE entity_id='Robot001'
 AND attribute_name='battery'
 ORDER BY time DESC LIMIT 20
 4. Returns time-series data
 ■■> Dashboard renders Chart.js line graph
 ■■> Shows battery level over time
-

Docker Network Architecture

All services run in Docker containers within the [fiware-network](#):

```
fiware-network (Bridge Network)
■
■■ fiware-mongo:27017
■  ■■> Backend for Orion Context Broker
■
■■ fiware-orion:1026
■  ■■> NGSI-LD API (public)
■  ■■> Connects to mongo-db
■
■■ telemetry-timescaledb:5432
■  ■■> PostgreSQL + TimescaleDB
■
■■ telemetry-service:8668
■  ■■> Flask API (public)
■  ■■> Connects to timescaledb
```

Port Mappings:

- `1026:1026` - Orion accessible from host
- `5432:5432` - TimescaleDB accessible from host (for admin)
- `8668:8668` - Telemetry Service accessible from host

Container Communication:

- Uses Docker DNS (container names as hostnames)
 - No need for localhost or IP addresses
 - Automatic service discovery
-

Technology Stack Summary

Component	Technology	Version	Purpose
Dashboard	Chart.js	1.1.0	Visualize historical data

Context Broker FIWARE Orion-LD 1.5.1 Real-time entity management
Database (Orion) MongoDB 4.4 Context Broker persistence
Time-series DB TimescaleDB latest-pg15 Historical data storage
Telemetry Service Python Flask 3.0.0+ Notification processing
CORS Support flask-cors 4.0.0+ Cross-origin requests
Dashboard HTML/CSS/JS - Web visualization
Charting Chart.js 3.9.1 Data visualization
Containerization Docker Compose - Service orchestration
Data Format NGSI-LD / JSON-LD - Semantic interoperability

Key Design Decisions

1. Why NGSI-LD?

- **Semantic Interoperability**: Linked Data ensures standardized vocabulary
- **FIWARE Ecosystem**: Integrates with many FIWARE components
- **Context Management**: Built for IoT and digital twin use cases
- **Subscription Model**: Event-driven architecture for real-time updates

2. Why Custom Telemetry Service (not QuantumLeap)?

- **Simplicity**: Lightweight Flask app easier to customize
- **Control**: Full control over data format and storage
- **Debugging**: Easier to troubleshoot and add logging
- **Learning**: Educational value in understanding the flow

3. Why TimescaleDB?

- **Performance**: 10-100x faster than PostgreSQL for time-series
- **SQL Compatible**: Use familiar SQL queries
- **Compression**: Automatic data compression saves space
- **Analytics**: Built-in functions for time-series analysis

4. Why Docker?

- **Isolation**: Each service runs in its own container
- **Portability**: Works on Windows, Linux, macOS
- **Networking**: Automatic DNS and service discovery
- **Reproducibility**: Same environment everywhere

System Requirements

Hardware:

- CPU: 4+ cores recommended
- RAM: 8GB minimum, 16GB recommended
- Disk: 10GB+ free space

Software:

- Docker Desktop with WSL2 (Windows) or Docker Engine (Linux)
- Python 3.12+ (for local development)

- Modern web browser (Chrome, Firefox, Edge)
- Git for version control

Network:

- Ports 1026, 5432, 8668, 8081 available
- Internet access for Docker image downloads

Future Integration

NVIDIA Omniverse (Isaac Sim)

Role: Physics-based robot simulation

Integration Plan:

- Install Isaac Sim
- Create Python connector script
- Read robot state from simulation
- Publish to Orion Context Broker via REST API
- Subscribe to action commands from RL policy

Example Code:

```
import requests
from omni.isaac.core import World

world = World()
robot = world.scene.get_object("robot")

while True:
    # Get robot state
    position = robot.get_world_pose()[0]
    velocity = robot.get_linear_velocity()

    # Update Context Broker
    requests.patch(
        "http://localhost:1026/v2/entities/Robot001/attrs",
        json={
            "position": {"value": position.tolist(), "type": "Array"},
            "velocity": {"value": velocity.tolist(), "type": "Array"}
        }
    )
```

Reinforcement Learning Policy

Role: AI agent learning optimal robot behavior

Integration Plan:

- Implement RL algorithm (PPO, SAC, DQN)
- Define observation space (from robot state)
- Define action space (motor commands)
- Read state from Context Broker
- Compute action using trained model
- Send action to Isaac Sim via Context Broker

Example Code:

```
import requests
import numpy as np
from stable_baselines3 import PPO
```

```

model = PPO.load("trained_model.zip")

while True:
    # Get current state
    response = requests.get(
        "http://localhost:1026/ngsi-ld/v1/entities/urn:ngsi-ld:Robot:001"
    )
    state = extract_observation(response.json())

    # Compute action
    action = model.predict(state)

    # Send action to simulation
    requests.patch(
        "http://localhost:1026/v2/entities/Robot001/attrs",
        json={"action": {"value": action.tolist(), "type": "Array"}}
    )

---

```

Troubleshooting Guide

Common Issues

1. Context Broker not responding

```

# Check if container is running
docker ps | grep fiware-orion

# View logs
docker logs fiware-orion --tail 50

# Restart services
docker-compose restart orion

```

2. Telemetry notifications not working

```

# Check subscription status
curl http://localhost:1026/v2/subscriptions

# Verify telemetry service is running
curl http://localhost:8668/health

# Check telemetry service logs
docker logs telemetry-service --tail 50

```

3. Dashboard not showing data

```

# Check CORS headers
curl -H "Origin: http://localhost:8081" http://localhost:1026/version

# Verify data exists
curl http://localhost:1026/ngsi-ld/v1/entities/urn:ngsi-ld:Robot:001

# Check browser console for errors (F12)

```

4. Database connection issues

```

# Check if TimescaleDB is running
docker ps | grep timescaledb

# Connect to database

```

```
docker exec -it telemetry-timescaledb psql -U postgres -d telemetry

# Check table contents
SELECT COUNT(*) FROM telemetry_data;
```

Performance Considerations

Context Broker:

- MongoDB indexes on entity IDs
- Connection pooling
- Query optimization with attribute filters

Telemetry Service:

- Batch inserts to reduce database round-trips
- Async processing for high-throughput scenarios
- Connection pool to TimescaleDB

TimescaleDB:

- Hypertable automatically partitions by time
- Compression policies for old data
- Appropriate chunk interval (default: 7 days)

Dashboard:

- Debounced refresh (5 seconds)
- Efficient Chart.js rendering
- Limits on historical data points (last 20)

Security Considerations

Current Implementation:

- No authentication on APIs (development only)
- CORS set to `*` (allows all origins)
- Database passwords in docker-compose (not encrypted)

Production Recommendations:

- Enable authentication on Orion (OAuth2, API keys)
- Use HTTPS/TLS for all communications
- Restrict CORS to specific origins
- Use Docker secrets for sensitive data
- Implement API rate limiting
- Add network segmentation
- Regular security updates

Monitoring & Maintenance

Health Checks:

```
# Orion
curl http://localhost:1026/version

# Telemetry Service
curl http://localhost:8668/health
```

```
# TimescaleDB
docker exec telemetry-timescaledb pg_isready -U postgres
```

Backup Strategy:

```
# Backup MongoDB (Orion data)
docker exec fiware-mongo mongodump --out /backup

# Backup TimescaleDB
docker exec telemetry-timescaledb pg_dump -U postgres telemetry > backup.sql
```

Log Monitoring:

```
# View all container logs
docker-compose logs -f

# Specific service
docker logs -f fiware-orion
docker logs -f telemetry-service
```

Conclusion

This Digital Twin architecture provides a **robust, scalable, and extensible foundation** for robot reinforcement learning in logistics applications.

Key Strengths:

- ■ Real-time state management with NGSI-LD
- ■ Historical data retention with time-series optimization
- ■ Event-driven architecture via subscriptions
- ■ Web-based monitoring and visualization
- ■ Containerized for easy deployment
- ■ Standards-based (FIWARE, NGSI-LD)

Ready for:

- Integration with NVIDIA Isaac Sim
- Deployment of RL training pipelines
- Scalability to multiple robots
- Extension with additional FIWARE components

Next Steps:

Integrate Isaac Sim simulation
Implement RL training loop
Add authentication and security
Scale to multiple robot entities
Deploy to production environment

Project Repository:

<https://github.com/fedih/Digital-Twin-NVIDIA-Omniverse-Robot-RL-for-Logistics>

Documentation Date: November 24, 2025